

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2023 NVIDIA Corporation
Built on Tue_Aug_15_22:02:13_PDT_2023
Cuda compilation tools, release 12.2, V12.2.140
Build cuda_12.2.r12.2/compiler.33191640_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-ei625o8p
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-ei625o8p
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit 5741c522547756ac4bb7a16df32106a15efb8a57
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
```

```
%load_ext nvcc4jupyter
```

```
Detected platform "Colab". Running its setup...
Source files will be saved in "/tmp/tmp6ar9c_mt".
```

```
%%cuda
```

```
#include<stdio.h>
#include<cuda.h>
#include<stdlib.h>
#include<time.h>
```

```
__global__ void max1(int* input)
{
    const int tid = threadIdx.x;

    auto step_size = 1;
    int number_of_threads = blockDim.x;
    int temp;

    while (number_of_threads > 0)
    {
        if (tid < number_of_threads) // still alive?
        {
            const auto fst = tid * step_size * 2;
            const auto snd = fst + step_size;
            //input[fst] += input[snd];
            if (input[fst]<input[snd])
            {
                temp=input[fst];
                input[fst]=input[snd];
                input[snd]=temp;
            }
        }
        __syncthreads();
        step_size <= 1;
        number_of_threads >>= 1;
    }
}
```

```
int main()
{
    const auto count = 8;
    const int size = count * sizeof(int);
    int h[] = {13, 65, 15, 14, 33, 2, 30, 8};

    int* d;

    cudaMalloc(&d, size);
    cudaMemcpy(d, h, size, cudaMemcpyHostToDevice);

    max1 <<<1, count / 2 >>>(d);

    int result;
    cudaMemcpy(&result, d, sizeof(int), cudaMemcpyDeviceToHost);
    // cout << "Large no is %d " << result << endl;
    printf("Large no is %d ", result);

    getchar();

    cudaFree(d);
    //delete[] h;
```

```
    return 0;  
}
```

Large no is 65

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2023 NVIDIA Corporation
Built on Tue_Aug_15_22:02:13_PDT_2023
Cuda compilation tools, release 12.2, V12.2.140
Build cuda_12.2.r12.2/compiler.33191640_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-ei625o8p
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-ei625o8p
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit 5741c522547756ac4bb7a16df32106a15efb8a57
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
```

```
%load_ext nvcc4jupyter
```

```
Detected platform "Colab". Running its setup...
Source files will be saved in "/tmp/tmp6ar9c_mt".
```

```
%%cuda
```

```
#include<stdio.h>
#include<cuda.h>
#include<stdlib.h>
#include<time.h>
```

```
__global__ void min1(int* input)
{
    const int tid = threadIdx.x;

    auto step_size = 1;
    int number_of_threads = blockDim.x;
    int temp;

    while (number_of_threads > 0)
    {
        if (tid < number_of_threads) // still alive?
        {
            const auto fst = tid * step_size * 2;
            const auto snd = fst + step_size;
            //input[fst] += input[snd];
            if (input[fst]>input[snd])
            {
                temp=input[fst];
                input[fst]=input[snd];
                input[snd]=temp;
            }
        }
        __syncthreads();
        step_size <= 1;
        number_of_threads >= 1;
    }
}
```

```
int main()
{
    const auto count = 8;
    const int size = count * sizeof(int);
    int h[] = {13, 65, 15, 14, 33, 23, 30, 8};

    int* d;

    cudaMalloc(&d, size);
    cudaMemcpy(d, h, size, cudaMemcpyHostToDevice);

    min1 <<<1, count / 2 >>>(d);

    int result;
    cudaMemcpy(&result, d, sizeof(int), cudaMemcpyDeviceToHost);
    // cout << "Large no is %d " << result << endl;
    printf("Small no is %d ", result);

    getchar();

    cudaFree(d);
    //delete[] h;
```

```
    return 0;  
}
```

Small no is 8

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2023 NVIDIA Corporation
Built on Tue_Aug_15_22:02:13_PDT_2023
Cuda compilation tools, release 12.2, V12.2.140
Build cuda_12.2.r12.2/compiler.33191640_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-ei625o8p
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-ei625o8p
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit 5741c522547756ac4bb7a16df32106a15efb8a57
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
```

```
%load_ext nvcc4jupyter
```

```
Detected platform "Colab". Running its setup...
Source files will be saved in "/tmp/tmp6ar9c_mt".
```

```
%%cuda
#include<stdio.h>
#include<cuda.h>
#include<stdlib.h>
#include<time.h>

__global__ void sum(int* input)
{
    const int tid = threadIdx.x;

    auto step_size = 1;
    int number_of_threads = blockDim.x;

    while (number_of_threads > 0)
    {
        if (tid < number_of_threads) // still alive?
        {
            const auto fst = tid * step_size * 2;
            const auto snd = fst + step_size;
            input[fst] += input[snd];
        }
        __syncthreads();
        step_size <= 1;
        number_of_threads >>= 1;
    }
}

int main()
{
    const auto count = 8;
    const int size = count * sizeof(int);
    int h[] = {13, 27, 15, 14, 33, 2, 30, 8};

    int* d;

    cudaMalloc(&d, size);
    cudaMemcpy(d, h, size, cudaMemcpyHostToDevice);

    sum <<<1, count / 2 >>>(d);

    int result;
    cudaMemcpy(&result, d, sizeof(int), cudaMemcpyDeviceToHost);
    // cout << "Sum is " << result << endl;
    printf("Sum is %d ", result);

    getchar();

    cudaFree(d);
    //delete[] h;

    return 0;
}
```

Sum is 142

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2023 NVIDIA Corporation
Built on Tue_Aug_15_22:02:13_PDT_2023
Cuda compilation tools, release 12.2, V12.2.140
Build cuda_12.2.r12.2/compiler.33191640_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-438dnets
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-438dnets
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit 5741c522547756ac4bb7a16df32106a15efb8a57
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Building wheels for collected packages: nvcc4jupyter
  Building wheel for nvcc4jupyter (pyproject.toml) ... done
  Created wheel for nvcc4jupyter: filename=nvcc4jupyter-1.2.1-py3-none-any.whl size=10739 sha256=652cfa192a926bb82952bb8d8fb3fd55f14
  Stored in directory: /tmp/pip-ephem-wheel-cache-2wx_nqnt/wheels/a8/b9/18/23f8ef71ceb0f63297dd1903aedd067e6243a68ea756d6feea
Successfully built nvcc4jupyter
Installing collected packages: nvcc4jupyter
Successfully installed nvcc4jupyter-1.2.1
```

```
%load_ext nvcc4jupyter
```

```
Detected platform "Colab". Running its setup...
Source files will be saved in "/tmp/tmp6jg70rkt".
```

```
%%cuda
/*[m,n]*[n,j]= [m,j]
first take transpose of vector anfd
then mulriply and add row wisw for each col.
*/
//parallelism of multiplication only
//tested ok on 05/12/2018

#include <cuda.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#define m 10
__global__ void mul_r(int *a, int *b, int *c){

    int tid = threadIdx.x;
    if (tid < m){
        c[tid]= a[tid] * b[tid];
    }

}

int main(){
    int n, c, d, fst[10][10], snd[10][10], t_snd[10][10];
    int row,col,sum_c, a[10], b[10], ans[10];

    /*
    printf("Enter the number of rows and columns of matrix\n");
    scanf("%d%d", &m, &n);
    */
    n=m; //square matrix only
    //printf("Enter the elements of first matrix\n");
    //for true random value of vector
    //srand(time(0));
    for (c = 0; c < m; c++)
    {

        for (d = 0; d < n; d++)
        {

            //scanf("%d", &first[c][d]);

            fst[c][d]=rand()%10+1;
        }
    }
    printf("display the elements of first matrix\n");
    for (c = 0; c < m; c++) {
        for (d = 0 ; d < n; d++) {
```

```

        printf("%d\t", fst[c][d]);
    }
    printf("\n");
}

// take next matrix
//for true random value of vector
//srand(time(0));
for (c = 0; c < m; c++)
{
    for (d = 0; d < n; d++)
    {
        //scanf("%d", &first[c][d]);

        snd[c][d]=rand()%10+1;
    }
}
printf("display the elements of second matrix\n");
for (c = 0; c < m; c++) {
    for (d = 0 ; d < n; d++) {

        printf("%d\t", snd[c][d]);
    }
    printf("\n");
}
// transpose of second matrix
for(c=0; c<m; c++)
    for(d=0; d<n; d++)
    {
        t_snd[d][c] = snd[c][d];
    }

// Displaying the transpose of matrix a
printf("\nTranspose of second Matrix:\n");
for (c = 0; c < n; c++) {
    for (d = 0 ; d < m; d++) {

        printf("%d\t", t_snd[c][d]);

    }
    printf("\n");
}

// now multiply on cuda
int *dev_a, *dev_b,*dev_ans;
cudaError_t err=cudaSuccess;
// allocate memory on GPU
err=cudaMalloc((void*)&dev_a,m * sizeof(int));
if (err !=cudaSuccess)
{ printf("failed to allocate on device \n");
printf("error is:\n",cudaGetErrorString(err));
exit(EXIT_FAILURE);
}
//printf("first ok\n");
cudaMalloc((void*)&dev_b,m * sizeof(int));
cudaMalloc((void*)&dev_ans,m * sizeof(int));
//printf("first finished ok\n");
row=0;
col=0;
cudaEvent_t start, end;
cudaEventCreate(&start);
cudaEventCreate(&end);
cudaEventRecord(start);

for(row=0; row<m; row++){

    for (d = 0 ; d < m; d++) {

        a[d]=fst[row][d];

    }
    // printf("ok a\n");
    cudaMemcpy(dev_a,a,m*sizeof(int), cudaMemcpyHostToDevice);
    for (col=0; col<m; col++){
        for (d= 0 ; d < m; d++) {

            b[d]=t_snd[col][d];
            ans[d]=0;

        }
        // printf("ok b\n");
        cudaMemcpy(dev_b,b,m*sizeof(int), cudaMemcpyHostToDevice);

```

```

    cudaMemcpy(dev_ans,ans,m*sizeof(int), cudaMemcpyHostToDevice);
//  printf("calling GPU\n");
    mul_r<<<1,m>>>(dev_a,dev_b,dev_ans);
    err=cudaMemcpy(ans,dev_ans,m*sizeof(int), cudaMemcpyDeviceToHost);
    if (err !=cudaSuccess)
    {   printf("failed to copy from device \n");
        exit(EXIT_FAILURE);
    }
//printf("GPU returned\n");
//a=fst[0];
sum_c=0;
for (d = 0 ; d < m; d++) {

    //printf("%d\t", ans[d]);
    sum_c+=ans[d];
}
snd[row][col]=sum_c;


//printf("one element=%d\n",snd[row][col]);
// printf("\n");

}
}
//
cudaEventRecord(end);
cudaEventSynchronize(end);
float time = 0;
cudaEventElapsedTime(&time, start, end);
printf("execution time=%f\n",time);
//
printf(" Matrix multipliation ans=:\n");
for (c = 0; c < n; c++) {
    for (d = 0 ; d < m; d++) {

        printf("%d\t", snd[c][d]);

    }
    printf("\n");
}

//
return 0;
}

display the elements of first matrix
4      7      8      6      4      6      7      3      10      2
3      8      1      10     4      7      1      7      3      7
2      9      8      10     3      1      3      4      8      6
10     3      3      9      10     8      4      7      2      3
10     4      2      10     5      8      9      5      6      1
4      7      2      1      7      4      3      1      7      2
6      6      5      8      7      6      7      10     4      8
5      6      3      6      5      8      5      5      4      1
8      9      7      9      9      5      4      2      5      10
3      1      7      9      10     3      7      7      5      10

display the elements of second matrix
6      1      5      9      8      2      8      3      8      3
3      7      2      1      7      2      6      10     5      10
1      10     2      8      8      2      2      6      10     8
8      7      8      4      7      6      7      4      10     5
9      2      3      10     4      10     1      9      9      6
1      10     7      4      9      6      7      2      2      6
10     9      5      9      2      1      4      1      5      5
5      5      8      7      4      2      8      6      10     7
3      2      8      9      6      8      5      2      9      6
10     8      6      4      9      9      4      2      9      10

Transpose of second Matrix:
6      3      1      8      9      1      10     5      3      10
1      7      10     7      2      10     9      5      2      8
5      2      2      8      3      7      5      8      8      6
9      1      8      4      10     4      9      7      9      4
8      7      8      7      4      9      2      4      6      9
2      2      2      6      10     6      1      2      8      9
8      6      2      7      1      7      4      8      5      4
3      10     6      4      9      2      1      6      2      2
8      5      10     10     9      2      5      10     9      9
3      10     8      5      6      6      5      7      6      10

execution time=3.148256
Matrix multipliation ans=:
278    357    303    377    361    261    288    251    428    372
290    323    301    264    348    268    300    248    389    355
289    342    287    316    358    263    274    268    451    385
353    323    330    400    375    295    327    276    456    348
350    344    352    406    369    266    349    236    439    340

```


| | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 198 | 196 | 186 | 254 | 238 | 205 | 183 | 196 | 274 | 253 |
| 404 | 413 | 374 | 427 | 418 | 319 | 359 | 306 | 526 | 445 |
| 249 | 295 | 265 | 301 | 303 | 218 | 269 | 223 | 341 | 301 |
| 405 | 408 | 342 | 418 | 463 | 360 | 336 | 329 | 535 | 463 |
| 413 | 381 | 345 | 429 | 379 | 345 | 287 | 272 | 525 | 412 |

Start coding or [generate](#) with AI.

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2023 NVIDIA Corporation
Built on Tue_Aug_15_22:02:13_PDT_2023
Cuda compilation tools, release 12.2, V12.2.140
Build cuda_12.2.r12.2/compiler.33191640_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-438dnets
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-438dnets
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit 5741c522547756ac4bb7a16df32106a15efb8a57
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Building wheels for collected packages: nvcc4jupyter
  Building wheel for nvcc4jupyter (pyproject.toml) ... done
  Created wheel for nvcc4jupyter: filename=nvcc4jupyter-1.2.1-py3-none-any.whl size=10739 sha256=652cfa192a926bb82952bb8fb3fd55f14
  Stored in directory: /tmp/pip-ephem-wheel-cache-2wx_nqnt/wheels/a8/b9/18/23f8ef71ceb0f63297dd1903aedd067e6243a68ea756d6feea
Successfully built nvcc4jupyter
Installing collected packages: nvcc4jupyter
Successfully installed nvcc4jupyter-1.2.1
```

```
%load_ext nvcc4jupyter
```

```
Detected platform "Colab". Running its setup...
Source files will be saved in "/tmp/tmp6jg70rkt".
```

```
%%cuda
#include<stdio.h>
#include<cuda.h>
#include<stdlib.h>
#include<time.h>

#define N 500

__global__ void add(int *a, int *b, int *c){

int tid = threadIdx.x;
if (tid < N){
    c[tid]= a[tid] + b[tid];
}

}

int main (void){
int a[N], b[N], c[N];
int *dev_a, *dev_b,*dev_c;
cudaError_t err=cudaSuccess;
err=cudaMalloc((void**)&dev_a,N * sizeof(int));
if (err !=cudaSuccess)
{ printf("failed to allocate on device \n");
printf("error is:\n",cudaGetErrorString(err));
exit(EXIT_FAILURE);
}
cudaMalloc((void**)&dev_b,N * sizeof(int));
cudaMalloc((void**)&dev_c,N * sizeof(int));

for(int i=0;i<N;i++){
a[i] =i;
b[i] = i*i;

c[i]=0;
}

/*for(int i=0;i<N;i++){
printf(" c contents are =%d\n", c[i]);
} */
cudaEvent_t start, end;
cudaEventCreate(&start);
cudaEventCreate(&end);
cudaEventRecord(start);

cudaMemcpy(dev_a,a,N*sizeof(int), cudaMemcpyHostToDevice);
cudaMemcpy(dev_b,b,N*sizeof(int), cudaMemcpyHostToDevice);
cudaMemcpy(dev_c,c,N*sizeof(int), cudaMemcpyHostToDevice);
add<<1,N>>>(dev_a,dev_b,dev_c);
```


```
err=cudaMemcpy(c,dev_c,N*sizeof(int), cudaMemcpyDeviceToHost);
if (err !=cudaSuccess)
{ printf("failed to copy from device \n");
exit(EXIT_FAILURE);
}

cudaEventRecord(end);
cudaEventSynchronize(end);
float time = 0;
cudaEventElapsedTime(&time, start, end);
printf("execution time=%f\n",time);

for(int i=0;i<N;i++){
//printf("%d +%d=%d\n", a[i], b[i], c[i]);
}

cudaFree(dev_a);
cudaFree(dev_b);
cudaFree(dev_c);

return 0;
}
```

 execution time=0.202368

Assignment 3

```
In [3]: from __future__ import absolute_import, division, print_function
# TensorFlow and tf.keras
import tensorflow as tf
from tensorflow import keras
# Helper libraries
import numpy as np
import matplotlib.pyplot as plt
```

```
In [4]: fashion_mnist = keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 [=====] - 6s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 [=====] - 1s 0us/step
```

```
In [5]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
```

```
In [6]: train_images.shape
```

```
Out[6]: (60000, 28, 28)
```

```
In [7]: len(train_labels)
```

```
Out[7]: 60000
```

```
In [8]: train_labels
```

```
Out[8]: array([9, 0, 0, ..., 3, 0, 5], dtype=uint8)
```


```
In [9]: test_images.shape
```

```
Out[9]: (10000, 28, 28)
```

```
In [10]: len(test_labels)
```

```
Out[10]: 10000
```

```
In [11]: plt.figure()
plt.imshow(train_images[0])
plt.colorbar()
plt.grid(False)
plt.show()
```



```
In [12]: train_images = train_images / 255.0  
test_images = test_images / 255.0
```

```
In [16]: plt.figure(figsize=(10,10))  
for i in range(25):  
    plt.subplot(5,5,i+1)  
    plt.xticks([])  
    plt.yticks([])  
    plt.grid(False)  
    plt.imshow(train_images[i], cmap=plt.cm.binary)  
    plt.xlabel(class_names[train_labels[i]])  
plt.show()
```



```
In [17]: model = keras.Sequential([
keras.layers.Flatten(input_shape=(28, 28)),
keras.layers.Dense(128, activation=tf.nn.relu),
keras.layers.Dense(10, activation=tf.nn.softmax)
])
```

```
In [18]: model.compile(optimizer='adam',
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
```

```
In [19]: model.fit(train_images, train_labels, epochs=5)
```

```
Epoch 1/5
1875/1875 [=====] - 2s 811us/step - loss: 0.4927 -
accuracy: 0.8278
Epoch 2/5
1875/1875 [=====] - 1s 758us/step - loss: 0.3689 -
accuracy: 0.8667
Epoch 3/5
1875/1875 [=====] - 2s 828us/step - loss: 0.3341 -
accuracy: 0.8782
Epoch 4/5
1875/1875 [=====] - 1s 747us/step - loss: 0.3102 -
accuracy: 0.8862
Epoch 5/5
1875/1875 [=====] - 2s 811us/step - loss: 0.2929 -
accuracy: 0.8919
```

Out[19]: <keras.src.callbacks.History at 0x17f7cfaf0>

```
In [20]: test_loss, test_acc = model.evaluate(test_images, test_labels)
print('Test accuracy:', test_acc)
```

```
313/313 [=====] - 0s 552us/step - loss: 0.3479 - accuracy: 0.8757
Test accuracy: 0.8756999969482422
```

```
In [21]: predictions = model.predict(test_images)
```

```
313/313 [=====] - 0s 468us/step
```

```
In [22]: predictions[0]
```

```
Out[22]: array([6.44061117e-07, 8.45806625e-09, 9.75436592e-07, 5.93774641e-10,
              8.83455755e-07, 1.17749525e-02, 1.09047915e-06, 6.85961824e-03,
              1.04943738e-05, 9.81351376e-01], dtype=float32)
```

```
In [23]: np.argmax(predictions[0])
```

```
Out[23]: 9
```

```
In [24]: test_labels[0]
```

```
Out[24]: 9
```

```
In [25]: def plot_image(i, predictions_array, true_label, img):
          predictions_array, true_label, img = predictions_array[i], true_label[i]
          plt.grid(False)
          plt.xticks([])
          plt.yticks([])

          plt.imshow(img, cmap=plt.cm.binary)

          predicted_label = np.argmax(predictions_array)
          if predicted_label == true_label:
              color = 'blue'
          else:
              color = 'red'

          plt.xlabel("{} {:2.0f}% ({})".format(class_names[predicted_label],
                                                100*np.max(predictions_array),
                                                class_names[true_label]),
                    color=color)

          def plot_value_array(i, predictions_array, true_label):
              predictions_array, true_label = predictions_array[i], true_label[i]
              plt.grid(False)
              plt.xticks([])
              plt.yticks([])
              thisplot = plt.bar(range(10), predictions_array, color="#777777")
              plt.ylim([0, 1])
              predicted_label = np.argmax(predictions_array)


              thisplot[predicted_label].set_color('red')
              thisplot[true_label].set_color('blue')
```

```
In [26]: i = 0
          plt.figure(figsize=(6,3))
          plt.subplot(1,2,1)
          plot_image(i, predictions, test_labels, test_images)
          plt.subplot(1,2,2)
```


```
plot_value_array(i, predictions, test_labels)
plt.show()
```




Ankle boot 98% (Ankle boot)



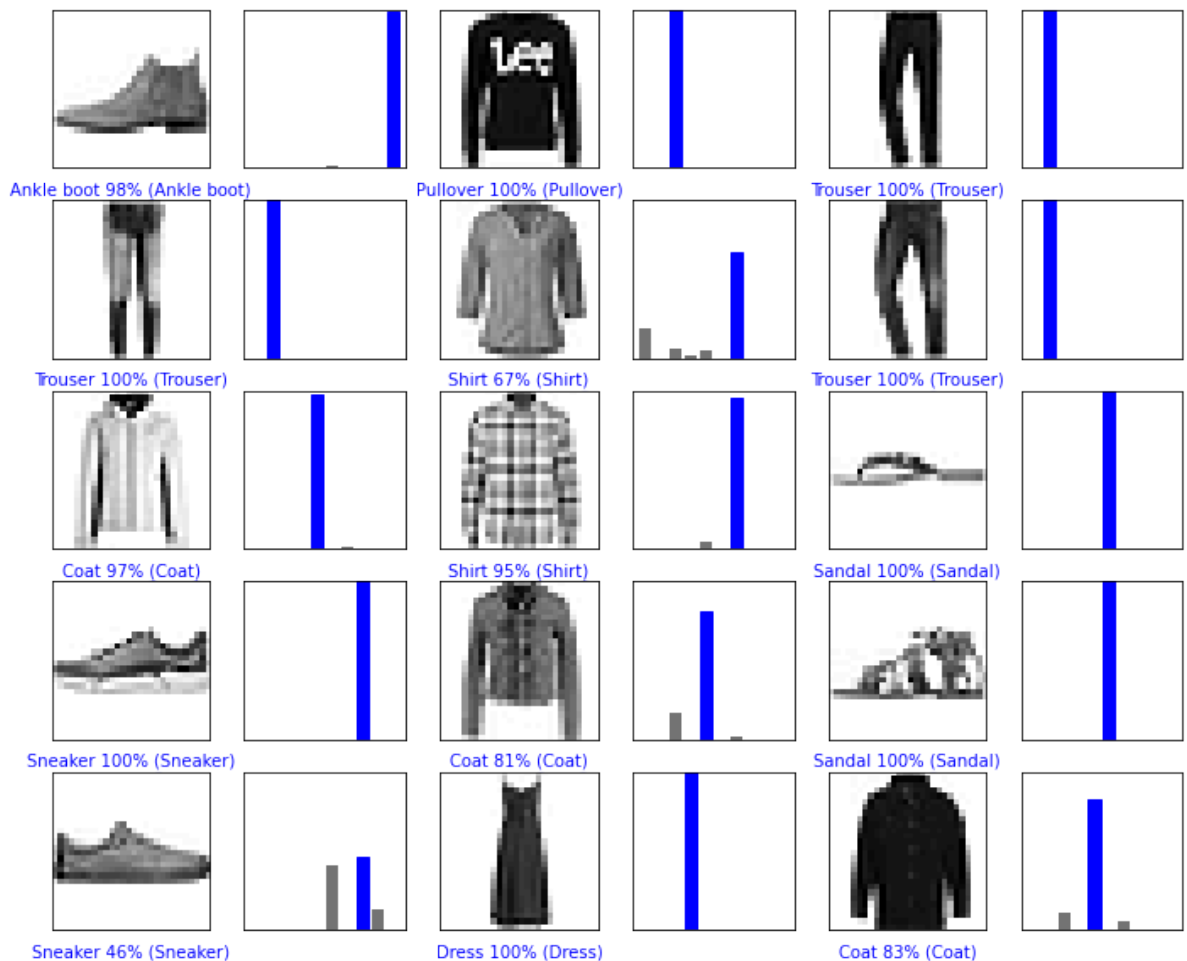
```
In [27]: i = 12
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plot_image(i, predictions, test_labels, test_images)
plt.subplot(1,2,2)
plot_value_array(i, predictions, test_labels)
plt.show()
```



Sneaker 46% (Sneaker)



```
In [28]: num_rows = 5
num_cols = 3
num_images = num_rows*num_cols
plt.figure(figsize=(2*2*num_cols, 2*num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2*num_cols, 2*i+1)
    plot_image(i, predictions, test_labels, test_images)
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_value_array(i, predictions, test_labels)
plt.show()
```

```
In [29]: img = test_images[0]
print(img.shape)

(28, 28)
```


```
In [30]: img = (np.expand_dims(img,0))
print(img.shape)

(1, 28, 28)
```

```
In [31]: predictions_single = model.predict(img)
print(predictions_single)

1/1 [=====] - 0s 13ms/step
[[6.4406112e-07 8.4580494e-09 9.7543830e-07 5.9377464e-10 8.8345740e-07
 1.1774947e-02 1.0904781e-06 6.8596182e-03 1.0494354e-05 9.8135138e-01]]
```

```
In [32]: plot_value_array(0, predictions_single, test_labels)
_ = plt.xticks(range(10), class_names, rotation=45)
```



```
In [33]: np.argmax(predictions_single[0])
```

```
Out[33]: 9
```